



# Deep Learning

# Content

---

- **Vanishing Gradient & Activation Functions**
- **Dropout**
- **Batch Normalization**




# Gradient Vanishing & Activation Functions

# Gradient Vanishing & Exploding

## ■ Gradient is easy to vanish or explode

- To many terms are multiplied.
- If some are small numbers, gradient becomes very small.
- If some are large numbers, gradient becomes very large.



$$\frac{\partial E}{\partial w_{ip}} = \left( \sum_{j=1}^J \left( \sum_{k=1}^K -(t_n - o_{nk}) o_{nk} (1 - o_{nk}) w_{kj} \right) h_{nj} (1 - h_{nj}) w_{ji} \right) h_{ni} (1 - h_{ni}) h_{np}$$



$$\frac{\partial E_n}{\partial w_{ji}} = -h_{nj} (1 - h_{nj}) x_{ni} \sum_{k=1}^m w_{kj} (t_{nk} - o_{nk}) o_{nk} (1 - o_{nk})$$

$$\frac{\partial E_n}{\partial w_{kj}} = -(t_{nk} - o_{nk}) o_{nk} (1 - o_{nk}) h_{nj}$$

# Activation Function

## ■ Vanishing Gradient

- The major terms are the derivatives of the activation function



$$\frac{\partial \text{Sigmoid}}{\partial w} \leq \frac{1}{4}$$


$$\frac{\partial E_n}{\partial w_{kj}} = -(t_{nk} - o_{nk}) o_{nk} (1 - o_{nk}) h_{nj}$$

$$\frac{\partial E_n}{\partial w_{ji}} = -h_{nj}(1 - h_{nj}) x_{ni} \sum_{k=1}^m w_{kj} (t_{nk} - o_{nk}) o_{nk} (1 - o_{nk})$$


$$\frac{\partial E}{\partial w_{ip}} = \left( \sum_{j=1}^J \left( \sum_{k=1}^K -(t_{nk} - o_{nk}) o_{nk} (1 - o_{nk}) w_{kj} \right) h_{nj} (1 - h_{nj}) w_{ji} \right) h_{ni} (1 - h_{ni}) h_{np}$$

# Activation Function

- Using another functions instead of sigmoid
  - Rectified Linear Unit (ReLU)



$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$



$$\frac{\partial E_n}{\partial w_{ji}} = \text{Some formular with 3 multiplications of } \frac{\partial f}{\partial w}$$
$$\frac{\partial E_n}{\partial w_{hg}} = \text{Some formular with 4 multiplications of } \frac{\partial f}{\partial w}$$

# Activation Function

---

## ■ Advantage

- No vanishing gradient problems.
  - Deep networks can be trained without pre-training
- Sparse activation
  - In a randomly initialized network, only about 50% of hidden units are activated
- Fast computation:
  - 6 times faster than sigmoid function

## ■ Disadvantage


- Knockout Problem

# Activation Function

- You may use another


- Leaky ReLU

$$f(x) = \begin{cases} x & \text{if } x > 0 \\ 0.01x & \text{otherwise} \end{cases}$$




- Swish (or SiLU-Sigmoid Linear Unit)

$$f(x) = \frac{x}{1 + e^{-x}}$$





# Other Activation Functions








# Regularization

# Overfitting


## ■ Overfitting



Underfit



Overfit




# Regularization

---

- **What is Regularization**
  - Introducing additional information to prevent over-fitting
- **Approaches**
  - Proper Learning: Early stopping
  - Proper Structure: Weight decay, Dropout, DropConnect, Stochastic pooling

# Early Stopping

- Split data into 3 groups



# Weight Decay

---

## ■ L1 Regularization

- Leading most weights very close to zero
- Choosing a small subset of most important inputs
- Resistant to noise in the inputs.

$$\tilde{E}(\mathbf{w}) = E(\mathbf{w}) + \frac{\lambda}{2} |\mathbf{w}|$$


## ■ L2 Regularization

- Penalizing peaky weights
- Encouraging to use all of its inputs a little rather than using only some of its inputs a lot.


$$\tilde{E}(\mathbf{w}) = E(\mathbf{w}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

# Weight Decay


- **Complex Structure vs Simple Structure**



Node  
Removal




Link  
Removal




# Weight Decay

- **Complex Structure vs Simple Structure**



Set  
many links  
to zero




$|w|$  is large  $\leftrightarrow$  NN is Complex

$|w|$  is small  $\leftrightarrow$  NN is Simple



# Weight Decay


- **Example: Separating green and red**



L2 regularization strengths of 0.01, 0.1, and 1

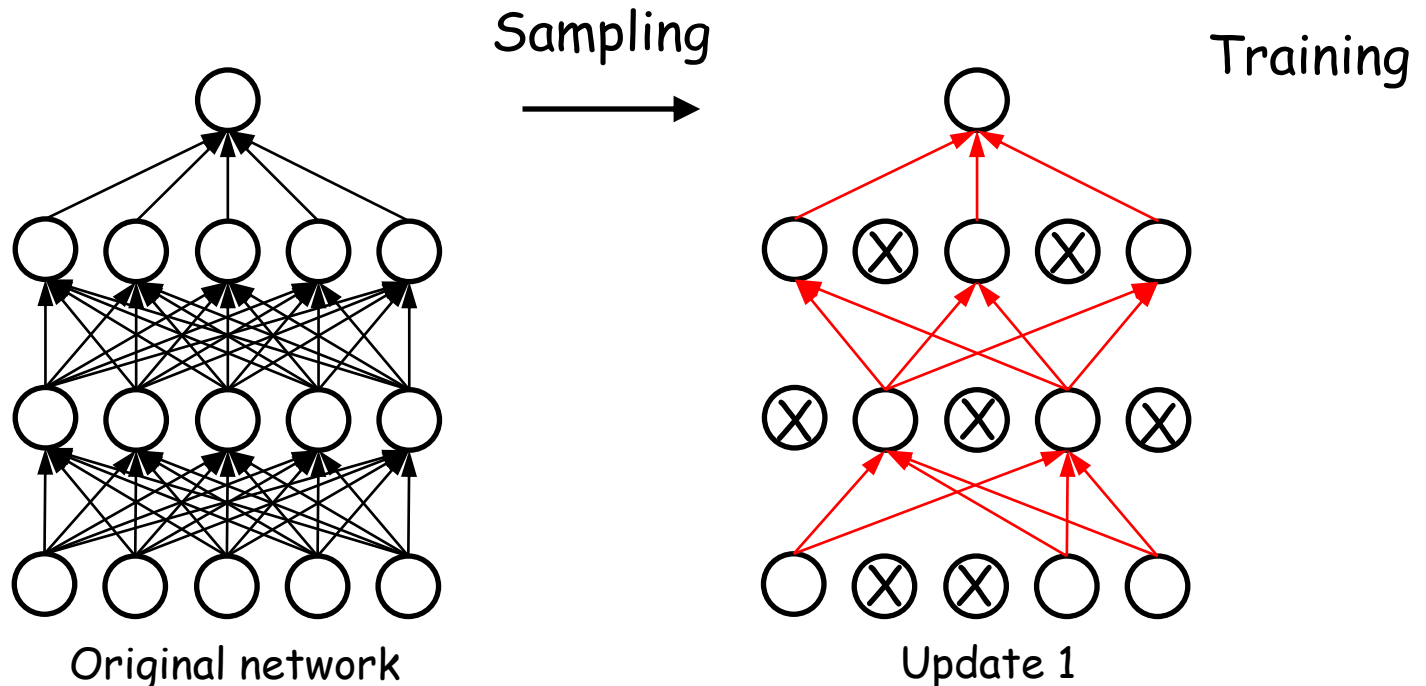
# Dropout

- **How can we reduce the structural complexity without removing nodes?**
  - Hmm??



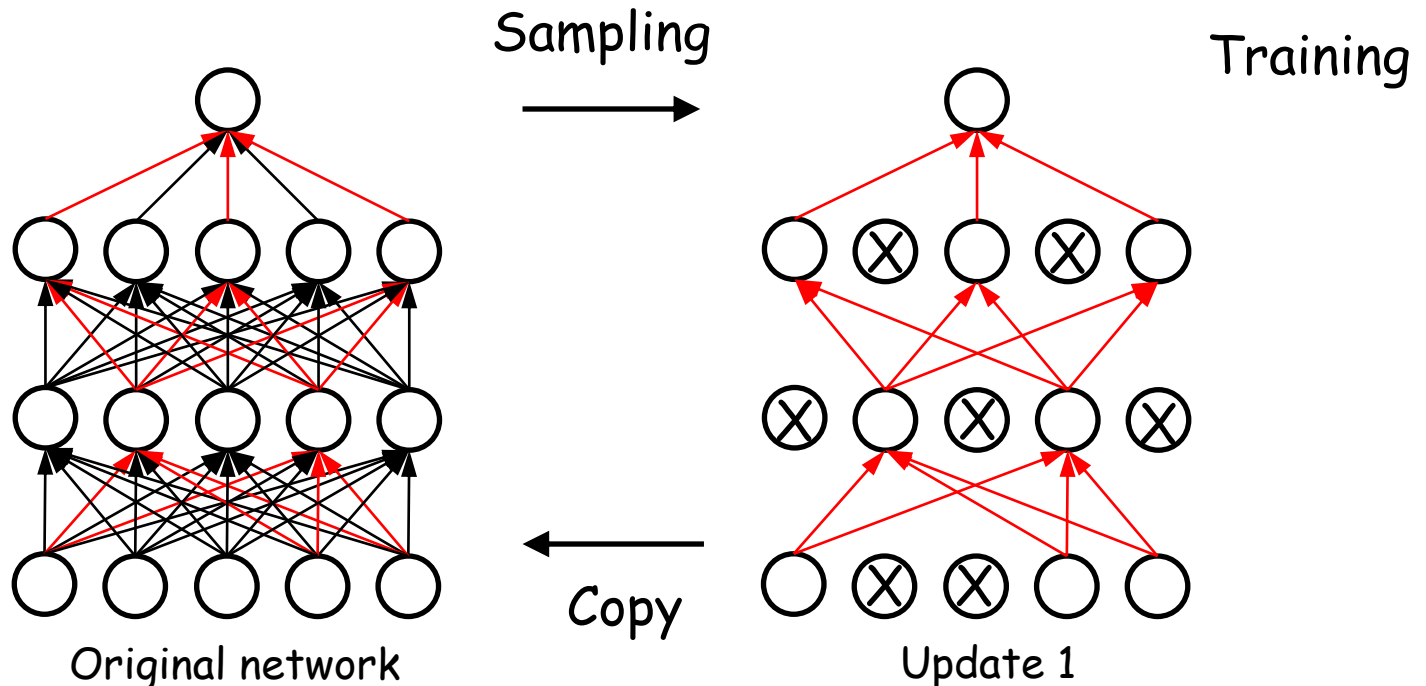
# Dropout

- **How can we reduce the structural complexity without removing nodes?**
  - Hmm??




# Dropout

- **How can we reduce the structural complexity without removing nodes?**
  - Hmm??




# Dropout

- **How can we reduce the structural complexity without removing nodes?**
  - Hmm??




# Dropout

- **How can we reduce the structural complexity without removing nodes?**
  - Hmm??



# Dropout


- **How can we reduce the structural complexity without removing nodes?**
  - Hmm??




# Dropout

- **Do this at every epoch**


- Randomly choose nodes with a probability of  $p$ 
  - Usually  $p = 0.5$
- Train the simplified neural network
  - At every epoch, we train different neural network which share connection weight each other




Original  
network



Update 1



Update 2




Update 3

...



# Dropout

- **Testing**
  - Use all the nodes without dropout




# Dropout

## ■ Testing


### Training of Dropout

Assume dropout rate is 50%



### Testing of Dropout


No dropout



# Dropout

## ■ The effect of the dropout rate $p$ :

- An architecture of 784-2048-2048-2048-10 is used on the MNIST dataset.
- The dropout rate  $p$  is changed from small numbers (most units are dropped out) to 1.0 (no dropout).



# Dropout

---

## ■ Summary



- Dropout is a very good and fast regularization method.
- Dropout is a bit slow to train (2-3 times slower than without dropout).
- If the amount of data is average-large – dropout excels. When data is big enough, dropout does not help much.
- Dropout achieves better results than former used regularization methods (Weight Decay).

# Batch Normalization


# Batch Normalization

- **Distribution Shift**


- Output distribution of the red node



Update 1




Update 2



Update 3



# Batch Normalization

- **Distribution Shift**
  - It disturbs the learning process,
  - Learning is getting slow down
- **Why don't we normalize the distribution of inputs**




# Batch Normalization


- Normalization of outputs




normalize ↪




Update 1




normalize ↪



Update 2



normalize ↪




Update 3





# Batch Normalization


## Input Normalization



$$\mu, \sigma^2 \quad \hat{net} = \frac{net - \mu}{\sqrt{\sigma^2 + \epsilon}}$$





Distorted distribution



Normalized distribution



# Batch Normalization

- Input Normalization



# Batch Normalization

## Input Normalization



$$net = \mathbf{w}\mathbf{x} + b$$


$$E(net) = E(\mathbf{w}\mathbf{x}) + b$$

$$\widehat{net} = \frac{\mathbf{w}\mathbf{x} + b - (E(\mathbf{w}\mathbf{x}) + b)}{\sqrt{\sigma^2 + \epsilon}}$$

$$\widehat{net} = \frac{\mathbf{w}\mathbf{x} - E(\mathbf{w}\mathbf{x})}{\sqrt{\sigma^2 + \epsilon}}$$

# Batch Normalization


## Input Normalization



$\mu, \sigma^2$



$$\widehat{net} = \frac{net - \mu}{\sqrt{\sigma^2 + \epsilon}}$$

$$\widetilde{net} = \gamma \widehat{net} + \beta$$



# Batch Normalization

- For a Single Node




$$\mu_B = \frac{1}{D_B} \sum_{d=1}^{D_B} net_{di}$$

$$\sigma_B^2 = \frac{1}{D_B} \sum_{d=1}^{D_B} (net_{di} - \mu)^2$$

# Batch Normalization


## ■ Testing

- For Training, the mean and variance of each batch are used for normalization
- For Testing, of which data the mean and variance will be used?
  - Estimated with those of batches in the training




# Batch Normalization


- Performance with BN



(a) *accuracy*



(b) Without BN



(c) With BN

*input distributions*

# Batch Normalization

---

## ■ Advantage

- Reduces internal covariant shift.
- Reduces the dependence of gradients on the scale of the connection weights.
- Regularizes the model and reduces the need for regularization techniques.
  - It adds some stochastic noise to the activations as a result of using noisy estimates computed on the mini-batches. This has a regularization effect in some applications,



# Batch Normalization

---

## ■ Disadvantage

- Expensive: Memory and time
  - Must keep interim results of all instances in a batch
  - Especially in CNN, usually an image is large
- Hard to apply when the batch size is small
  - If batches are small, the means and variances cannot approximate the global ones.
- Hard to apply to recurrent networks
  - It doesn't match to structure of recurrent networks
  - Hard to implement with recurrent networks